

David I. Atkinson

Email: davatk@gmail.com
GitHub: diatkinson
Scholar: David Atkinson
Phone: (240) 899-3524

<i>Research Interests</i>	Mechanistic interpretability, AI governance, language model introspection	
<i>Education</i>	Northeastern University (Boston, MA) Ph.D. in Computer Science Advisors: David Bau. GPA: 4.0/4.0.	September 2023 ↓ Present
	University of Colorado, Boulder (Boulder, CO) B.A. in Computer Science and Mathematics Advisor: Chenhao Tan. GPA: 3.9/4.0. Graduated with distinction.	January 2016 ↓ May 2019
	Deep Springs College (Deep Springs, CA) Coursework in philosophy and mathematics. Two-year junior college.	June 2011 ↓ June 2013
<i>Awards</i>	NSF Computer and Information Science and Engineering Graduate Fellow	September 2023
<i>Experience</i>	Epoch AI (Remote) Staff Researcher Conducted research on forecasting AI progress and capabilities, developing quantitative methods to model and predict AI's trajectory.	November 2022 ↓ July 2023
	Farmers Business Network (San Carlos, CA) Software Engineer Rewrote data pipelines for satellite imagery ingestion and cloud detection, improving response time and saving costs. Transitioned the authorization system to MemoryDB-based caches, reducing p99 latency by a factor of over 100. Improved the company's AWS SQS- and DynamoDB-based message queue, providing visibility into messages in flight and reducing average latency by 50% for certain workflows.	February 2020 ↓ October 2022
<i>Preprints</i>	Agents of Chaos Natalie Shapira, Chris Wendler, Avery Yen, Gabriele Sarti, Koyena Pal, Olivia Floody, Adam Belfki, Alex Loftus, Aditya Ratan Jannali, Nikhil Prakash, Jasmine Cui, Giordano Rogers, Jannik Brinkmann, Can Rager, Amir Zur, Michael Ripa, Aruna Sankaranarayanan, David Atkinson, Rohit Gandikota, Jaden Fiotto-Kaufman, EunJeong Hwang, Hadas Orgad, P Sam Sahil, Negev Taglicht, Tomer Shabtay, Atai Ambus, Nitay Alon, Shiri Oron, Ayelet Gordon-Tapiero, Yotam Kaplan, Vered Shwartz, Tamar Rott Shaham, Christoph Riedl, Reuth Mirsky, Maarten Sap, David Manheim, Tomer Ullman, David Bau	2026
<i>Publications</i>	Identifying Introspection from the Inside David Atkinson, Dillon Plunkett, David Bau <i>Conference on Language Modeling (CoLM) 2026</i>	2026

	Algorithmic progress in language models Anson Ho, Tamay Besiroglu, Ege Erdil, David Owen, Robi Rahman, Zifan C Guo, David Atkinson, Neil Thompson, Jaime Sevilla <i>Advances in Neural Information Processing Systems</i> , Volume 37, Pages 58245-58283	2024
	Token Erasure as a Footprint of Implicit Vocabulary Items in LLMs Sheridan Feucht, David Atkinson, Byron Wallace, David Bau <i>Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP)</i>	2024
	Locating and Editing Factual Associations in Mamba Arnab Sen Sharma, David Atkinson, David Bau <i>Conference on Language Modeling (CoLM) 2024</i>	2024
	What Gets Echoed? Understanding the ‘Pointers’ in Explanations of Persuasive Arguments David Atkinson, Kumar Bhargav Srinivasan, Chenhao Tan <i>Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing (EMNLP)</i>	2019
Workshops	Government AI Use as a Monitoring Primitive: A Public Document Pilot Study David Atkinson, Joan Eleanor O’Bryan <i>Second Workshop on Technical AI Governance Research at ICML 2026</i>	2026
	Position: Strategic Resilience Requires the Targeted Maintenance of Human Technical Capacities Joan Eleanor O’Bryan, David Atkinson <i>Second Workshop on Technical AI Governance Research at ICML 2026</i>	2026
	Testing Language Model Agents Safely in the Wild Silen Naihin, David Atkinson, Marc Green, Merwane Hamadi, Craig Swift, Douglas Schonholtz, Adam Tauman Kalai, David Bau <i>Socially Responsible Language Modelling Research (SoLaR) Workshop at NeurIPS 2023</i>	2023
Reports	Trading Off Compute in Training and Inference Pablo Villalobos, David Atkinson <i>Epoch AI</i>	2023
	Direct Approach Interactive Model David Atkinson, Matthew Barnett, Edu Roldán, Ben Cottier, Tamay Besiroglu <i>Epoch AI</i>	2023
	Literature Review of Transformative Artificial Intelligence Timelines Keith Wynroe, David Atkinson, Jaime Sevilla <i>Epoch AI</i>	2023
Invited Reviewing	Conference on Neural Information Processing Systems (NeurIPS) 2026 (Sydney, Australia)	2026
	Conference on Language Modeling (COLM) 2026 (San Francisco, CA)	2026

	International Conference on Machine Learning (ICML) 2026, Second Workshop on Technical AI Governance Research (Seoul, South Korea)	2026
	Conference on Neural Information Processing Systems (NeurIPS) 2025, Mechanistic Interpretability Workshop (San Diego, CA)	2025
	Conference on Language Modeling (COLM) 2025, Workshop on the Application of LLM Explainability to Reasoning and Planning (Montreal, QC)	2025
	Conference on Neural Information Processing Systems (NeurIPS) 2024, Workshop on Interpretable AI: Past, Present and Future (Vancouver, BC)	2024
<i>Etc.</i>	ARENA 3.0 (Remote) Intensive bootcamp covering technical AI safety and alignment research topics.	January 2024 ↓ February 2024
	BlueDot Impact (Remote) Facilitator for the AI Safety Fundamentals governance course. Led discussions on potential AI regulatory regimes and policy approaches.	July 2023 ↓ November 2023
	Intro to ML Safety, Center for AI Safety (Remote) Advanced track course on empirical topics in machine learning safety. Focus areas: robustness, interpretability, and alignment.	September 2022 ↓ November 2022
	Giving What We Can Pledged to give 10% of income to effective charities.	2016 ↓ Present
	Recurse Center (New York, NY) 3-month writer's retreat for programmers. Focused on type theory, functional programming with Haskell, and React Native.	September 2015 ↓ December 2015